

Pragma Analytics Software Suite “PASS”

Time series data solution

matthieu@pragma-innovation.fr

Pragma Innovation

- Company Positioning: IT services focused on a « practice »
 - Valuate timeseries oriented data,
 - Analyze, report, decision support tool,
 - Market: ISP, Hosting providers, Telco
- Model: integration of open source avec customization
 - Follows big data reference model,
 - Evaluate, choose, design ideal open source software,
 - Consultancy,
- 20 years of expertise and track of records
 - Complex solution design,
 - Consulting during key transition phases

IP networking big data use cases

Traffic engineering

- Peering analysis,
- Backbone engineering,
- Capacity planning,
- Latency checking,
- Network Troubleshooting and monitoring,
- Network automation,
- ...

Security

- Forensic and post mortem,
- Traffic anomaly,
 - DoS and DDoS,
- Remediation strategy
- Security Automation,
- ...

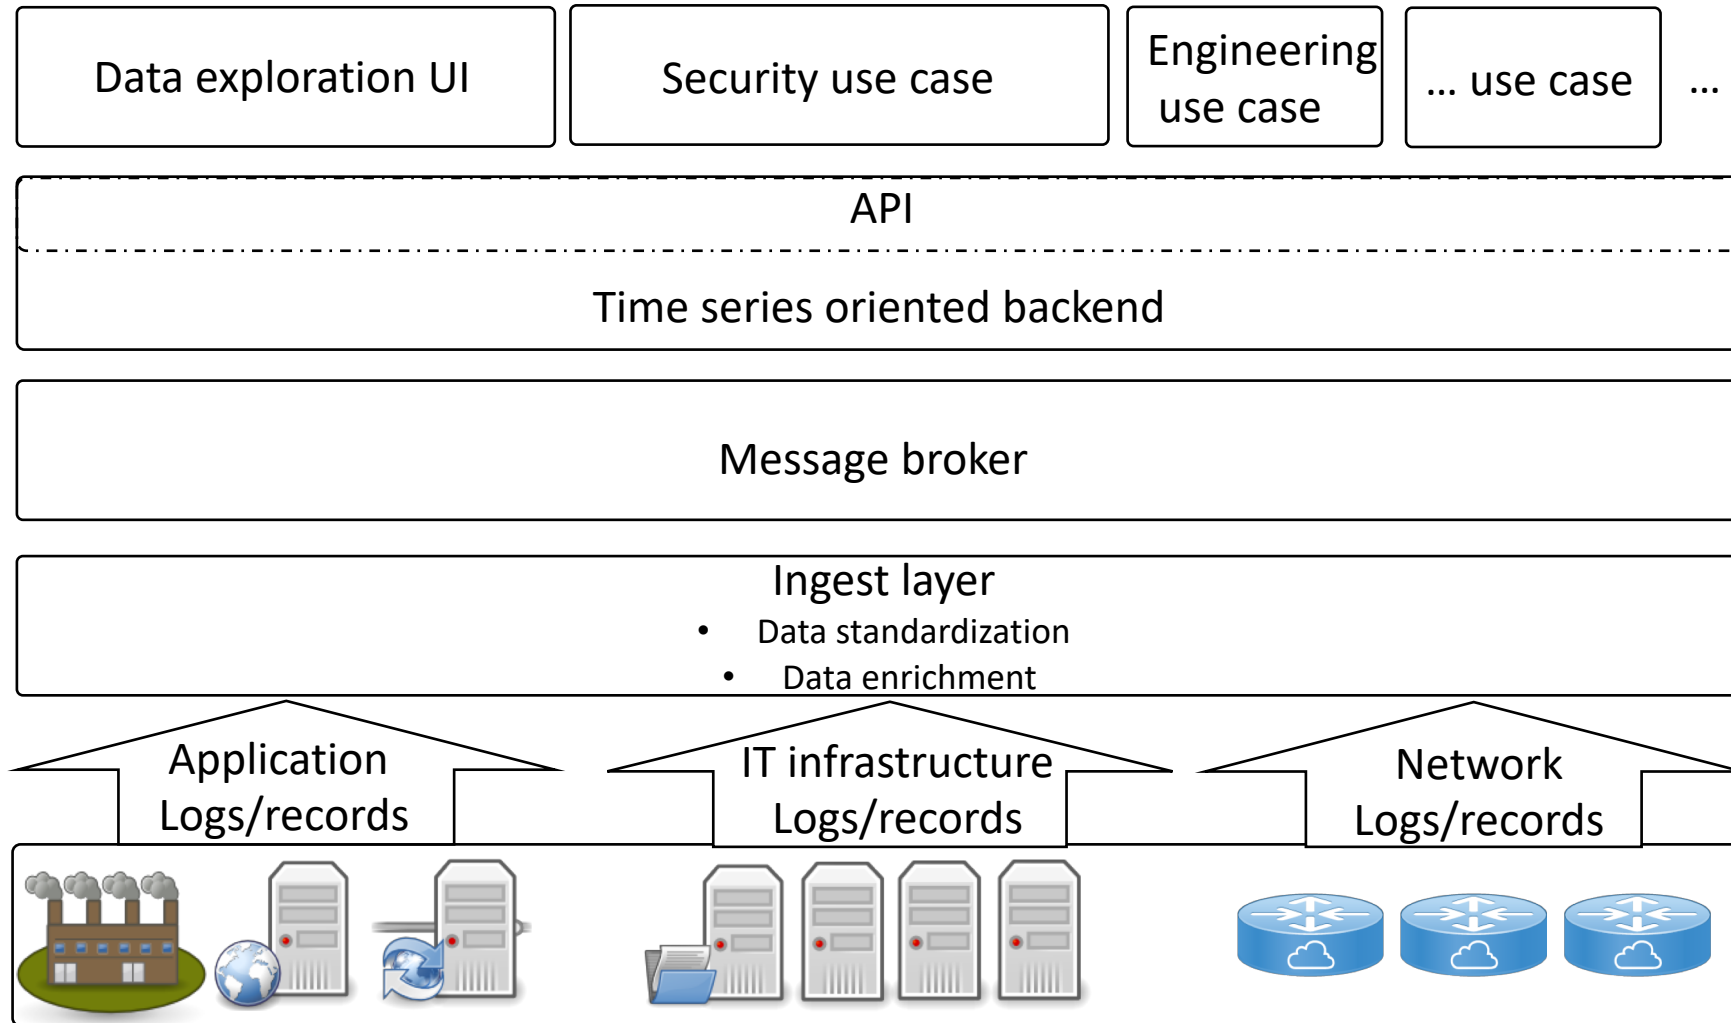
Business Intelligence

- Geo-marketing,
- Cost modeling,
- Profiling,
- Data mining,
- Service usage,
- OTT usage
- ...

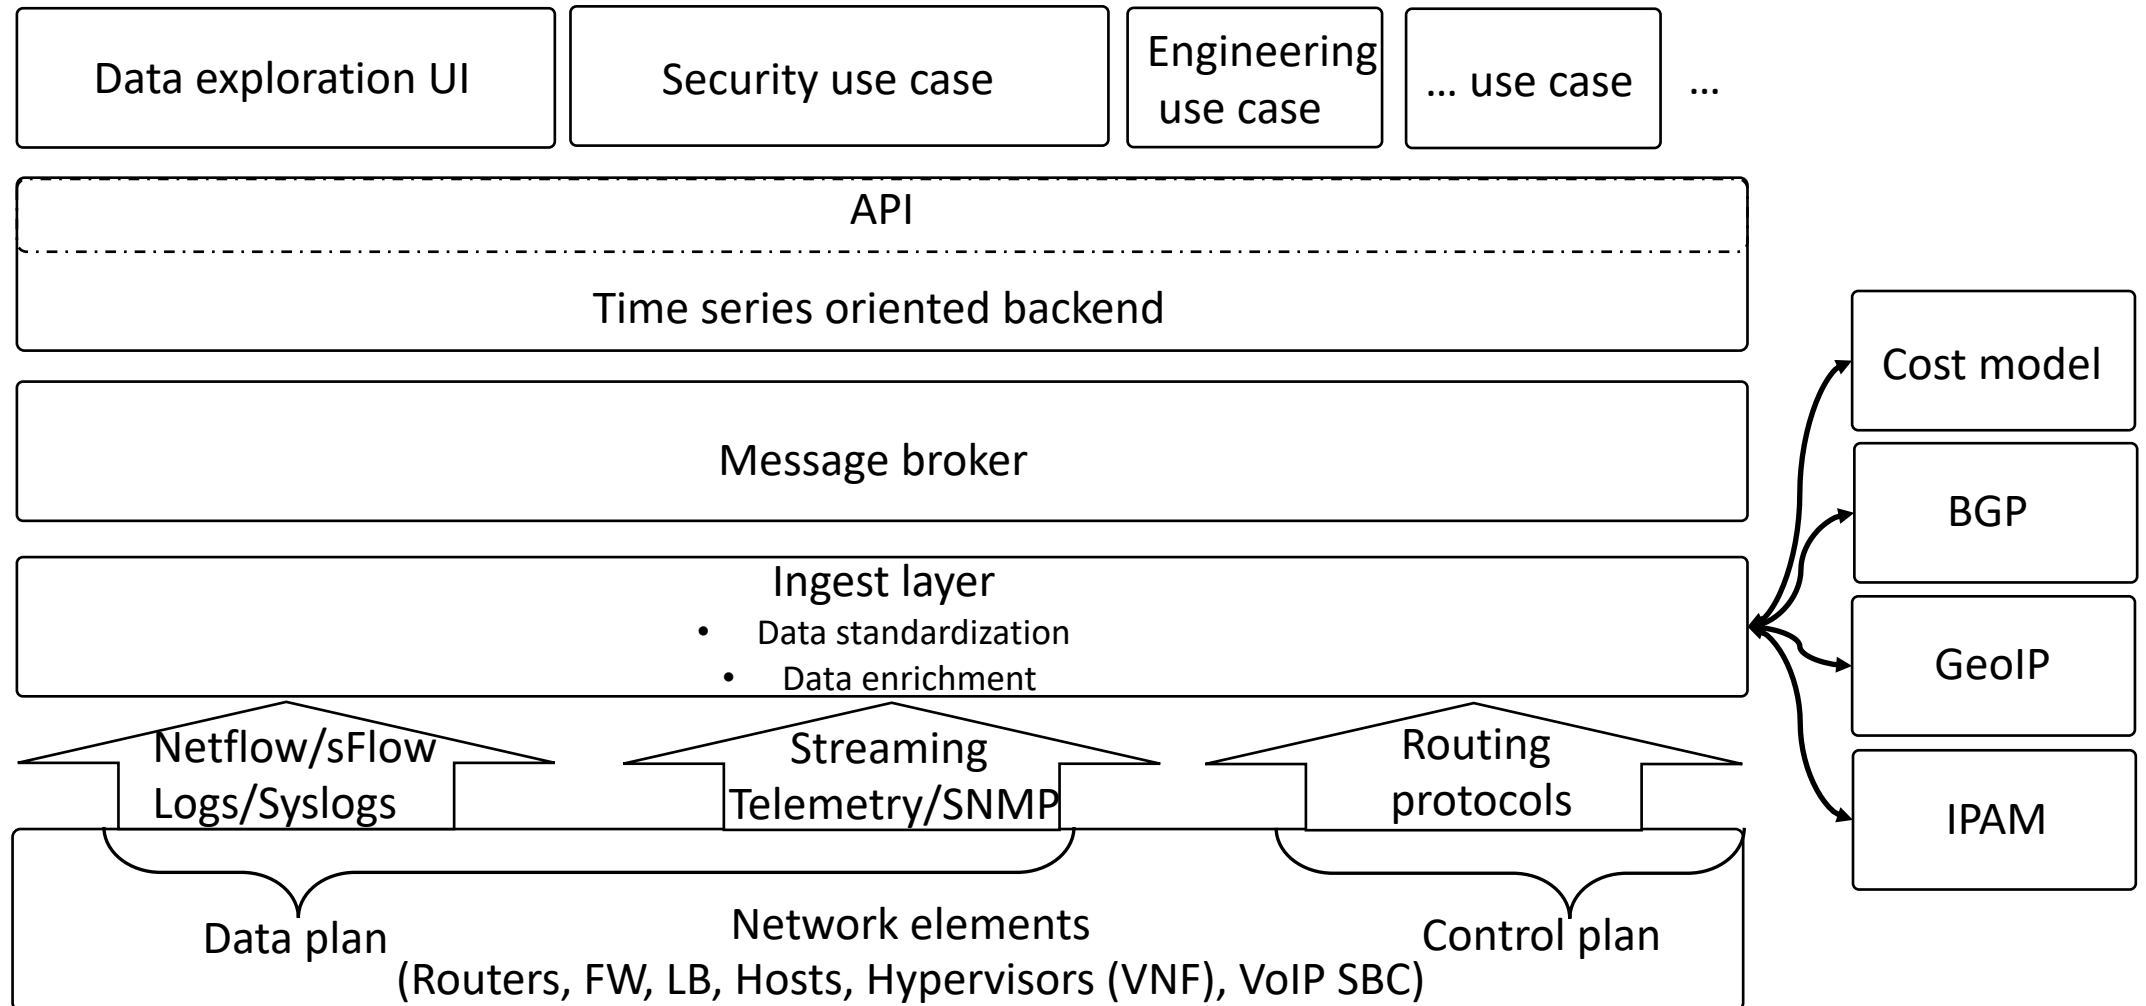
Compelling events for a game change

- Availability of open source time series backend solution
- Ingest tools of network data (open source)
- Community around open source UI analytics
- Appetite for valuable and customer specific big data solution
 - Companies data is part of their assets
- Need around customer specific use cases like
 - Cost optimization and network design as close as possible to enterprise specifics,
 - Application oriented SLA's and metrics,
 - Network process automation,

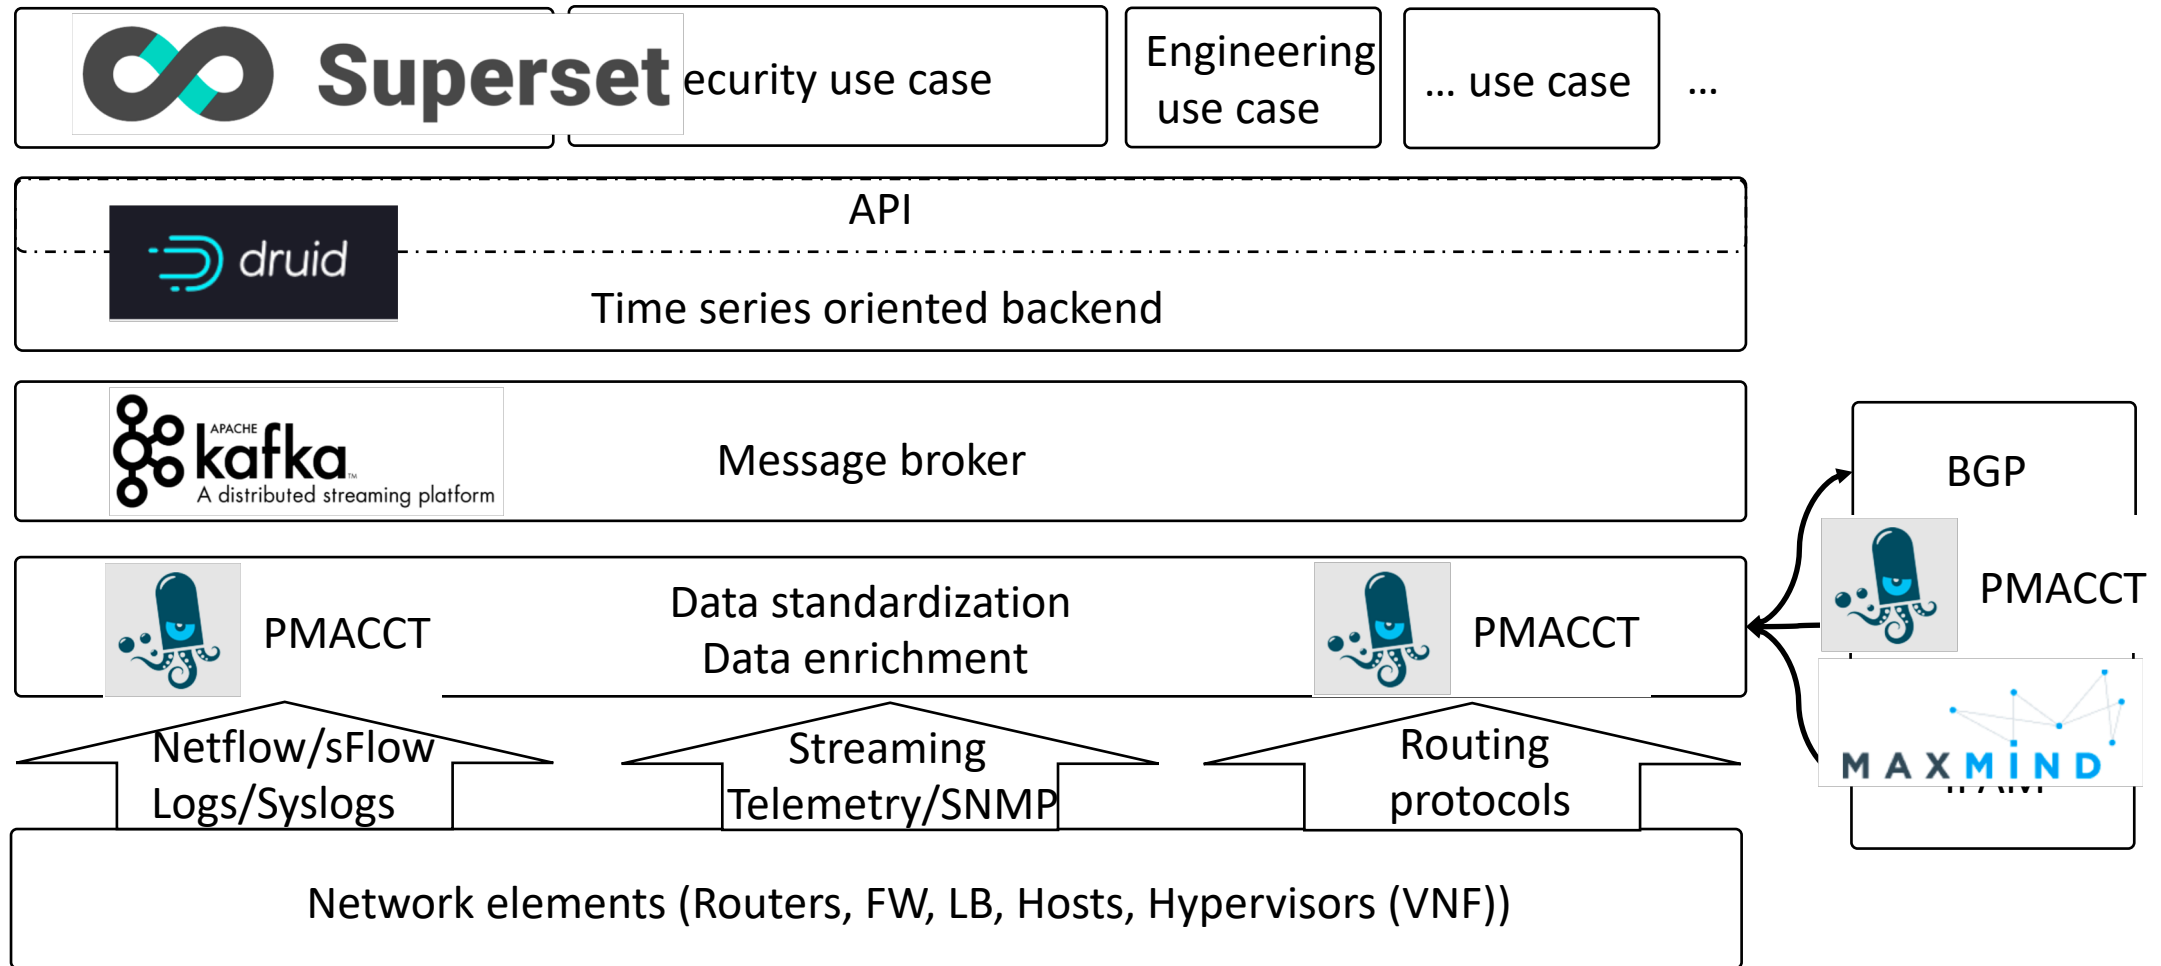
Data pipe line: reference model



Pipe line overview: Network telemetry



Pipe line overview: IP telemetry R1.0

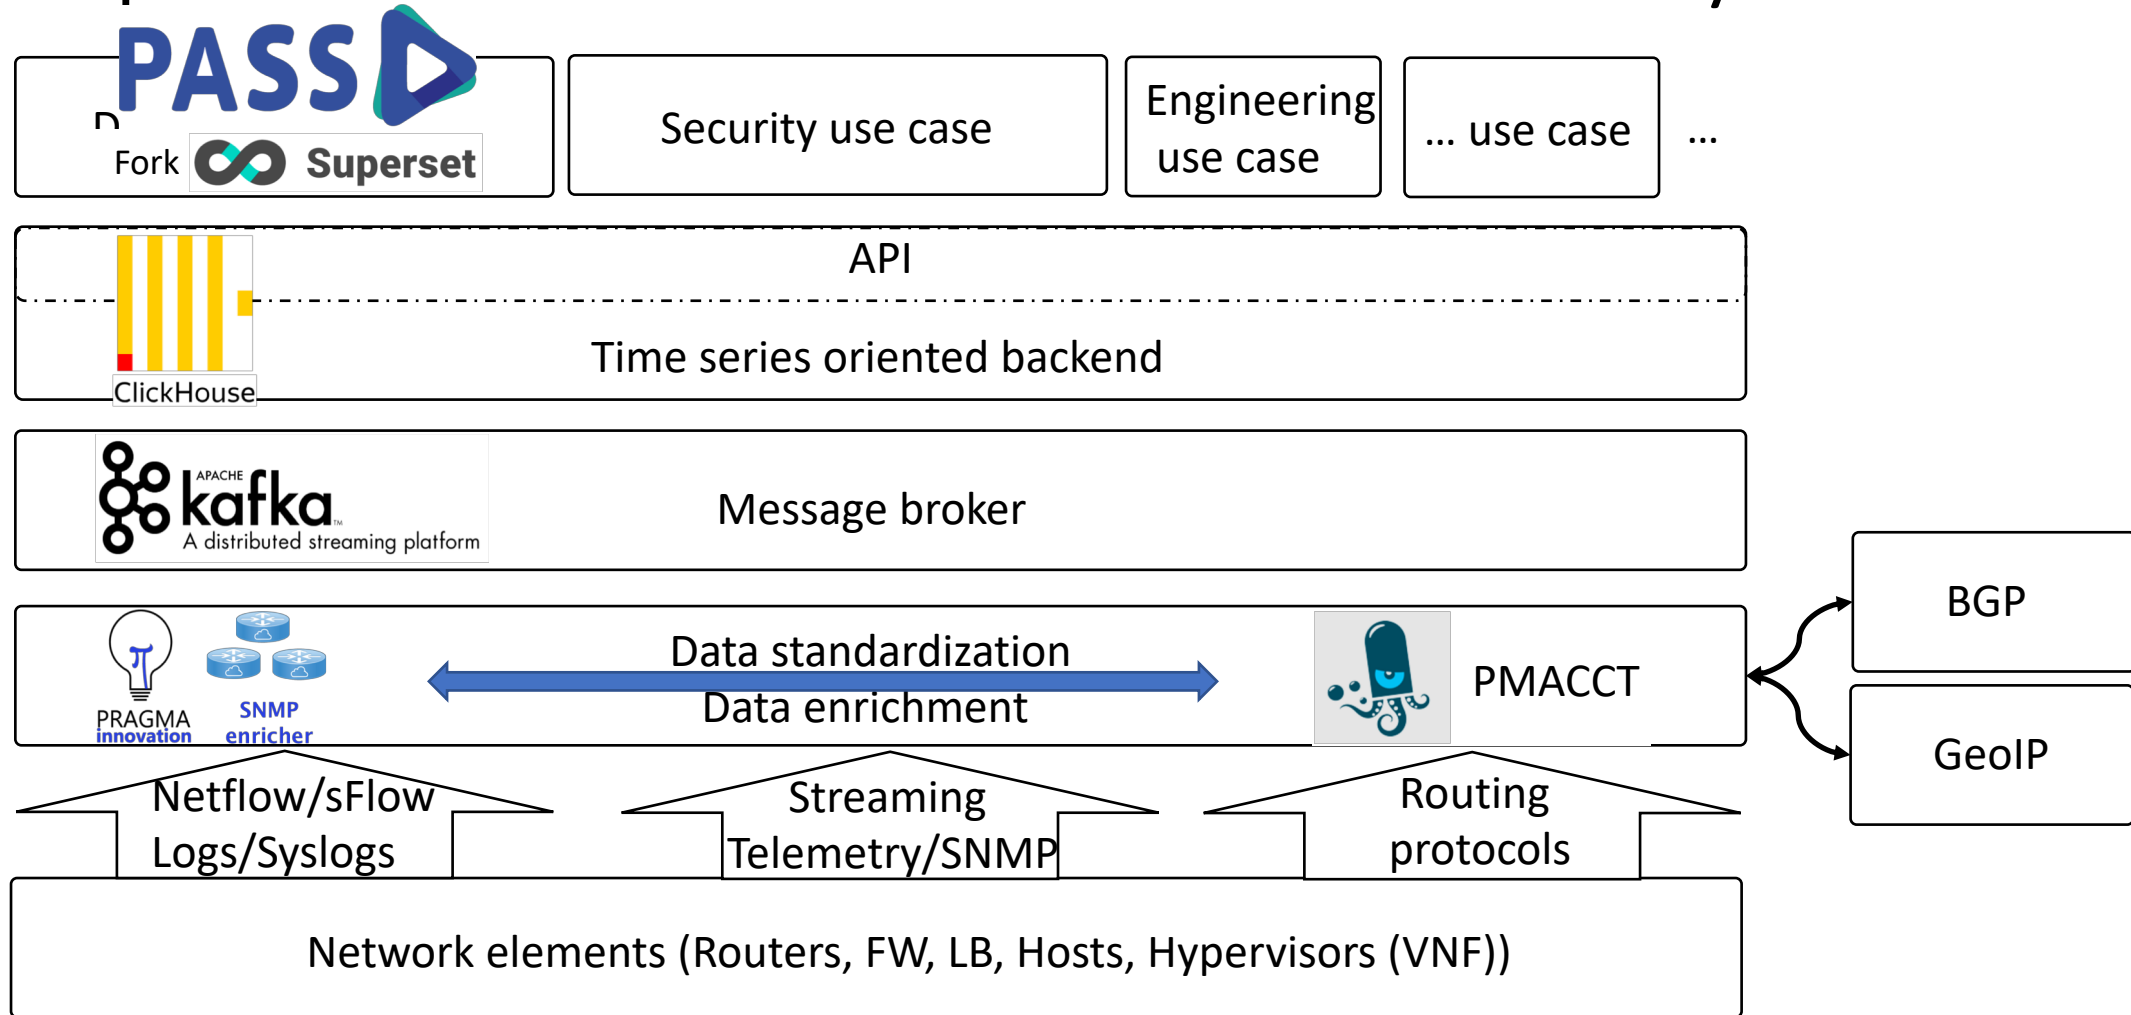


UI to be consolidated based on evaluation

R1.0 Main Issues highlighted with R 1.0

- Time granularity and data resampling:
 - We have to manage network bin cycle and user variable time granularity
 - Network bin cycle: SUM of flows that match the query
 - If User bin cycle > network bin cycle: User bin cycle: MAX/PERCENTILE of SUM(flows),
 - Not doable within Druid backend
- SNMP data enrichment : two options (not exclusive)
 - Create SNMP kafka streams,
 - Enrich Netflow/sFlow data,
- IPv4/IPv6 fields and CIDR block queries
 - Today IP address stored as string
 - Druid doesn't support of binary storage of IPv4 and IPv6
 - Druid doesn't support of CIDR block indexing
 - Blocking issue to enable CIDR filters or CIDR flow grouping
- Tech/Eco issue:
 - Druid cluster requires many nodes: zookeeper, meta data, druid data nodes, deep storage (HDFS), and an army of JVM's.
- Build a release 2.0 that address those issues

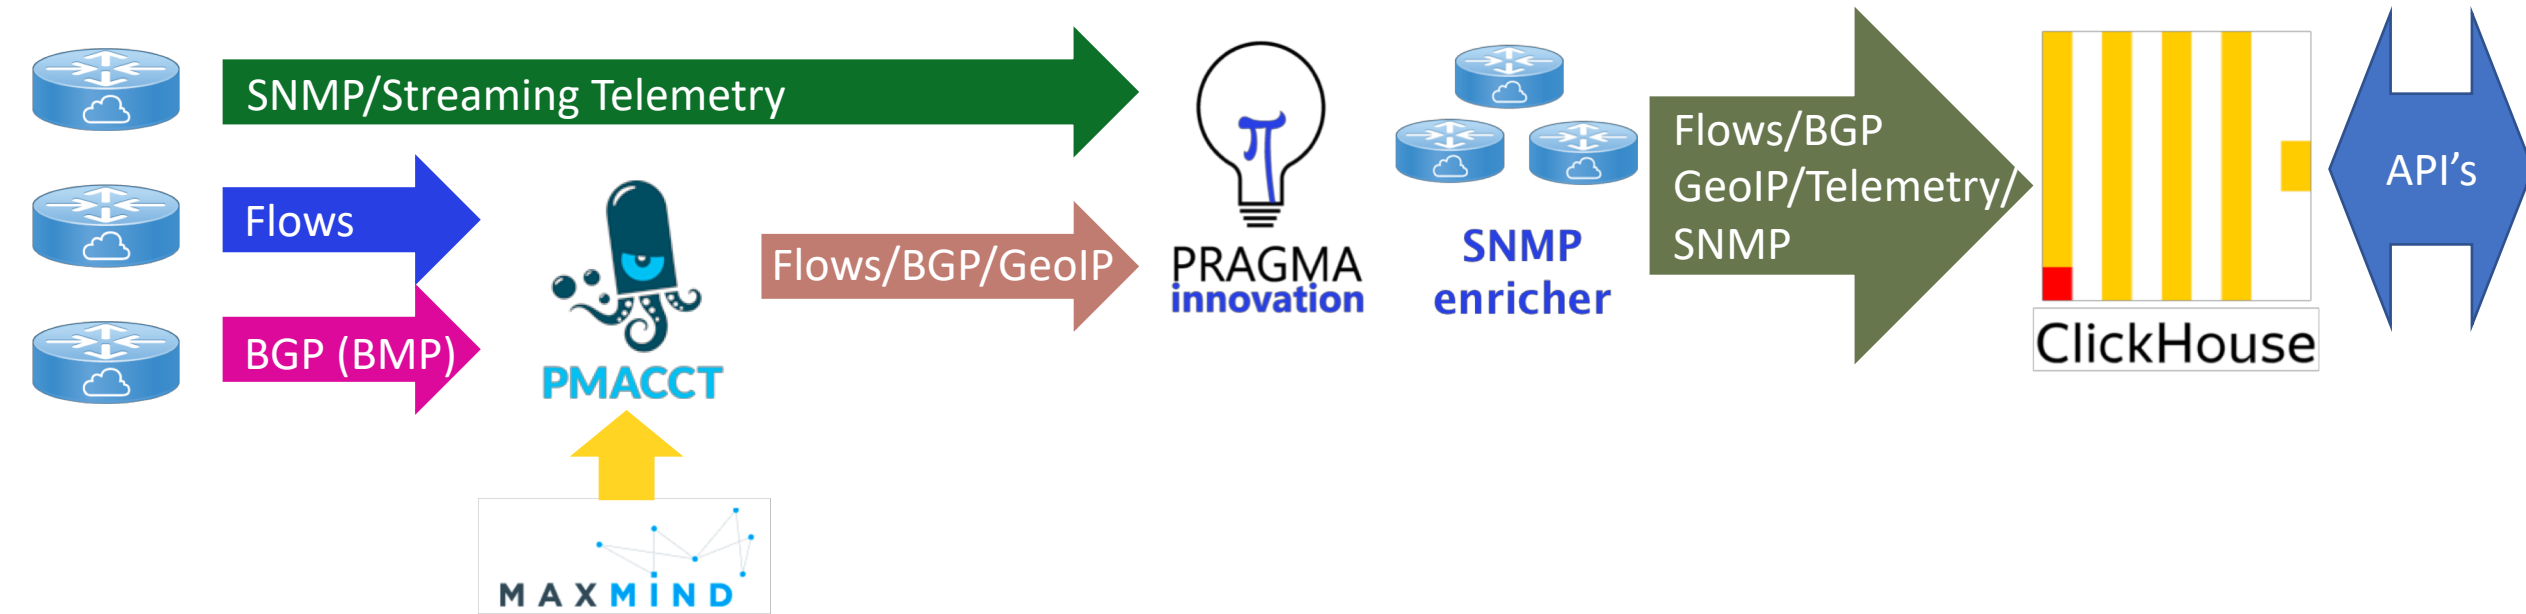
Pipe line overview: IP flow telemetry R2.0



Issues of 1.0 addressed by Clickhouse

- No need of complex node hierarchy and deep storage like Hadoop HDFS
 - Clickhouse manage data and data resiliency by itself no need to add other level of complexity
- Clickhouse support of IPv4 and IPv6 fields as well as operators of those fields,
- Data sampling and re-sampling : clickhouse is not "touching" or aggregating native metrics coming from network equipment and that's much much better that way !
 - Even better, it can do it if you want using Aggregating Merge tree engine

IP telemetry: Network flows pipeline




Message broker

- Kafka
 - Kafka support on PMACCT, SNMP streamer, SNMP enricher
 - Performances:
 - Kafka: 100K's events per second,
 - RabbitMQ: 20K+ events per second,
 - RabbitMQ more routing capabilities between large amount of process
 - Kafka more oriented for large amount of messages between simple entities,
- Value for Netflow/SNMP telemetry use case:
 - Netflow could be very bursty depending on implementation (due to active timeout, hardware implementation), Kafka will be very useful to absorb traffic spikes,
 - Kafka could help during backend upgrades (storing data during the upgrade operations)

Backend

- Clickhouse
 - Performance at ingest, performance to deal with queries,
 - Manage storage and resiliency without additional software (no HDFS),
- Value for the IP telemetry use case:
 - Clickhouse propose valuable features for Kafka ingest
 - Queuing, Views and calculation between Kafka JSON and table model,
 - IP telemetry metrics are about speed (PPS and BPS)
 - Clickhouse do support all mathematical calculation (MAX or SUM, Quantile, ...)
 - Clickhouse doesn't apply arbitrary metrics aggregation and keep it raw
- Open source (Yandex)
 - <https://github.com/yandex/ClickHouse>
 - CPP code with a decent amount of comments and specs
 - Good documentation quality: <https://clickhouse.yandex/docs/en/>
 - Company to provide advanced services Altinity

Example of design (Lab)



Access data and display dashboards

- Choice we made about frontend and UI is to use Superset with addition of custom modules
 - <https://github.com/apache/incubator-superset>
- Clickhouse data are reachable via SQLAlchemy and makes clickhouse working with all frameworks using SQLAlchemy
 - <https://github.com/xzkostyan/clickhouse-sqlalchemy>
 - Allows both access via HTTP and native interface
- Clickhouse and SQLAlchemy allow the use of PANDAS (lib python for “big data”)

Ingest data

- Native support of Kafka consumer
- Flexible implementation with some form of “SPARK like” features
 - Create a queue that map your data model (like JSON key/value)
 - Create your destination table which could differ from the queue table
 - Create a MATERIALIZED view between the two: you have all SQL tools to play with fields between the input queue table and the destination table
- Example with SNMP data:
 - Queue:

```
CREATE TABLE default.snmp_queue ( rtr_name String, if_name String, if_description String, if_speed UInt32, if_index UInt16, observation_time UInt32, date_calculation UInt64, in_octets_qty UInt64, in_unicast_pkts_qty UInt64, in_multicast_pkts_qty UInt64, in_broadcast_pkts_qty UInt64, out_octets_qty UInt64, out_unicast_pkts_qty UInt64, out_multicast_pkts_qty UInt64, out_broadcast_pkts_qty UInt64) ENGINE = Kafka('172.16.1.11:9092', 'snmp-streamer', 'clickhouse', 'JSONEachRow')
```


Ingest data

- Example with SNMP data:
 - Destination table:

```
CREATE TABLE default.snmp ( rtr_name String, if_name String, if_description String,  
if_speed UInt32, if_index UInt16, observation_time UInt32, date_calculation DateTime,  
day Date, in_octets_qty UInt64, in_unicast_pkts_qty UInt64, in_multicast_pkts_qty  
UInt64, in_broadcast_pkts_qty UInt64, out_octets_qty UInt64, out_unicast_pkts_qty  
UInt64, out_multicast_pkts_qty UInt64, out_broadcast_pkts_qty UInt64) ENGINE =  
MergeTree(day, (rtr_name, if_index, date_calculation), 8192)
```

Ingest data

- Example with SNMP data:
 - Materialized view playing with UNIX timestamp:

```
CREATE MATERIALIZED VIEW default.snmp_consumer TO default.snmp ( day Date,  
date_calculation DateTime, rtr_name String, if_name String, if_description String,  
if_speed UInt32, if_index UInt16, observation_time UInt32, in_octets_qty UInt64,  
in_unicast_pkts_qty UInt64, in_multicast_pkts_qty UInt64, in_broadcast_pkts_qty  
UInt64, out_octets_qty UInt64, out_unicast_pkts_qty UInt64, out_multicast_pkts_qty  
UInt64, out_broadcast_pkts_qty UInt64) AS SELECT toDate(date_calculation) AS day,  
toDateTime(date_calculation) AS date_calculation, rtr_name, if_name, if_description,  
if_speed, if_index, observation_time, in_octets_qty, in_unicast_pkts_qty,  
in_multicast_pkts_qty, in_broadcast_pkts_qty, out_octets_qty, out_unicast_pkts_qty,  
out_multicast_pkts_qty, out_broadcast_pkts_qty FROM default.snmp_queue
```

Thanks